

6.824 2018 Lecture 1: Introduction

6.824: Distributed Systems Engineering

What is a distributed system?

- multiple cooperating computers
- storage for big web sites, MapReduce, peer-to-peer sharing, &c
- lots of critical infrastructure is distributed

Why distributed?

- to organize physically separate entities
- to achieve security via isolation
- to tolerate faults via replication
- to scale up throughput via parallel CPUs/mem/disk/net

But:

- complex: many concurrent parts
- must cope with partial failure
- tricky to realize performance potential

Why take this course?

- interesting -- hard problems, powerful solutions
- used by real systems -- driven by the rise of big Web sites
- active research area -- lots of progress + big unsolved problems
- hands-on -- you'll build serious systems in the labs

COURSE STRUCTURE

<http://pdos.csail.mit.edu/6.824>

Course staff:

- Malte Schwarzkopf, lecturer
- Robert Morris, lecturer
- Deepti Raghavan, TA
- Edward Park, TA
- Erik Nguyen, TA
- Anish Athalye, TA

Course components:

- lectures
- readings
- two exams
- labs
- final project (optional)
- TA office hours
- piazza for announcements and lab help

Lectures:

- big ideas, paper discussion, and labs

Readings:

- research papers, some classic, some new
- the papers illustrate key ideas and important details
- many lectures focus on the papers
- please read papers before class!
- each paper has a short question for you to answer
- and you must send us a question you have about the paper
- submit question&answer by midnight the night before

Exams:

- Mid-term exam in class
- Final exam during finals week

Lab goals:

- deeper understanding of some important techniques
- experience with distributed programming
- first lab is due a week from Friday
- one per week after that for a while

Lab 1: MapReduce
Lab 2: replication for fault-tolerance using Raft
Lab 3: fault-tolerant key/value store
Lab 4: sharded key/value store

Optional final project at the end, in groups of 2 or 3.
The final project substitutes for Lab 4.
You think of a project and clear it with us.
Code, short write-up, short demo on last day.

Lab grades depend on how many test cases you pass
we give you the tests, so you know whether you'll do well
careful: if it often passes, but sometimes fails,
chances are it will fail when we run it

Debugging the labs can be time-consuming
start early
come to TA office hours
ask questions on Piazza

MAIN TOPICS

This is a course about infrastructure, to be used by applications.
About abstractions that hide distribution from applications.
Three big kinds of abstraction:
Storage.
Communication.
Computation.
[diagram: users, application servers, storage servers]

A couple of topics come up repeatedly.

Topic: implementation
RPC, threads, concurrency control.

Topic: performance
The dream: scalable throughput.
Nx servers -> Nx total throughput via parallel CPU, disk, net.
So handling more load only requires buying more computers.
Scaling gets harder as N grows:
Load im-balance, stragglers.
Non-parallelizable code: initialization, interaction.
Bottlenecks from shared resources, e.g. network.
Note that some performance problems aren't easily attacked by scaling
e.g. decreasing response time for a single user request
might require programmer effort rather than just more computers

Topic: fault tolerance
1000s of servers, complex net -> always something broken
We'd like to hide these failures from the application.
We often want:
Availability -- app can make progress despite failures
Durability -- app will come back to life when failures are repaired
Big idea: replicated servers.
If one server crashes, client can proceed using the other(s).

Topic: consistency
General-purpose infrastructure needs well-defined behavior.
E.g. "Get(k) yields the value from the most recent Put(k,v)."
Achieving good behavior is hard!
"Replica" servers are hard to keep identical.
Clients may crash midway through multi-step update.
Servers crash at awkward moments, e.g. after executing but before replying.
Network may make live servers look dead; risk of "split brain".
Consistency and performance are enemies.
Consistency requires communication, e.g. to get latest Put().
"Strong consistency" often leads to slow systems.
High performance often imposes "weak consistency" on applications.
People have pursued many design points in this spectrum.

CASE STUDY: MapReduce

Let's talk about MapReduce (MR) as a case study

MR is a good illustration of 6.824's main topics
and is the focus of Lab 1

MapReduce overview

context: multi-hour computations on multi-terabyte data-sets
e.g. analysis of graph structure of crawled web pages
only practical with 1000s of computers
often not developed by distributed systems experts
distribution can be very painful, e.g. coping with failure
overall goal: non-specialist programmers can easily split
data processing over many servers with reasonable efficiency.
programmer defines Map and Reduce functions
sequential code; often fairly simple
MR runs the functions on 1000s of machines with huge inputs
and hides details of distribution

Abstract view of MapReduce

input is divided into M files
[diagram: maps generate rows of K-V pairs, reduces consume columns]
Input1 -> Map -> a,1 b,1 c,1
Input2 -> Map -> b,1
Input3 -> Map -> a,1 c,1
 | | |
 | | -> Reduce -> c,2
 | -----> Reduce -> b,2
 -----> Reduce -> a,2

MR calls Map() for each input file, produces set of k2,v2
"intermediate" data
each Map() call is a "task"
MR gathers all intermediate v2's for a given k2,
and passes them to a Reduce call
final output is set of <k2,v3> pairs from Reduce()
stored in R output files
[diagram: MapReduce API --
map(k1, v1) -> list(k2, v2)
reduce(k2, list(v2) -> list(k2, v3)]

Example: word count

input is thousands of text files
Map(k, v)
 split v into words
 for each word w
 emit(w, "1")
Reduce(k, v)
 emit(len(v))

MapReduce hides many painful details:

starting s/w on servers
tracking which tasks are done
data movement
recovering from failures

MapReduce scales well:

N computers gets you Nx throughput.
Assuming M and R are $\geq N$ (i.e., lots of input files and map output keys).
Maps()s can run in parallel, since they don't interact.
Same for Reduce()s.
The only interaction is via the "shuffle" in between maps and reduces.
So you can get more throughput by buying more computers.
Rather than special-purpose efficient parallelizations of each application.
Computers are cheaper than programmers!

What will likely limit the performance?

We care since that's the thing to optimize.
CPU? memory? disk? network?

In 2004 authors were limited by "network cross-section bandwidth".

[diagram: servers, tree of network switches]

Note all data goes over network, during Map->Reduce shuffle.

Paper's root switch: 100 to 200 gigabits/second

1800 machines, so 55 megabits/second/machine.

Small, e.g. much less than disk (~50-100 MB/s at the time) or RAM speed.

So they cared about minimizing movement of data over the network.

(Datacenter networks are much faster today.)

More details (paper's Figure 1):

master: gives tasks to workers; remembers where intermediate output is

M Map tasks, R Reduce tasks

input stored in GFS, 3 copies of each Map input file

all computers run both GFS and MR workers

many more input tasks than workers

master gives a Map task to each worker

hands out new tasks as old ones finish

Map worker hashes intermediate keys into R partitions, on local disk

Q: What's a good data structure for implementing this?

no Reduce calls until all Maps are finished

master tells Reducers to fetch intermediate data partitions from Map workers

Reduce workers write final output to GFS (one file per Reduce task)

How does detailed design reduce effect of slow network?

Map input is read from GFS replica on local disk, not over network.

Intermediate data goes over network just once.

Map worker writes to local disk, not GFS.

Intermediate data partitioned into files holding many keys.

Q: Why not stream the records to the reducer (via TCP) as they are being produced by the mappers?

How do they get good load balance?

Critical to scaling -- bad for N-1 servers to wait for 1 to finish.

But some tasks likely take longer than others.

[diagram: packing variable-length tasks into workers]

Solution: many more tasks than workers.

Master hands out new tasks to workers who finish previous tasks.

So no task is so big it dominates completion time (hopefully).

So faster servers do more work than slower ones, finish abt the same time.

What about fault tolerance?

I.e. what if a server crashes during a MR job?

Hiding failures is a huge part of ease of programming!

Q: Why not re-start the whole job from the beginning?

MR re-runs just the failed Map(s) and Reduce(s).

MR requires them to be pure functions:

they don't keep state across calls,

they don't read or write files other than expected MR inputs/outputs,

there's no hidden communication among tasks.

So re-execution yields the same output.

The requirement for pure functions is a major limitation of

MR compared to other parallel programming schemes.

But it's critical to MR's simplicity.

Details of worker crash recovery:

* Map worker crashes:

master sees worker no longer responds to pings

crashed worker's intermediate Map output is lost

but is likely needed by every Reduce task!

master re-runs, spreads tasks over other GFS replicas of input.

some Reduce workers may already have read failed worker's intermediate data.

here we depend on functional and deterministic Map()!

master need not re-run Map if Reduces have fetched all intermediate data

though then a Reduce crash would then force re-execution of failed Map

* Reduce worker crashes.

finished tasks are OK -- stored in GFS, with replicas.

master re-starts worker's unfinished tasks on other workers.

* Reduce worker crashes in the middle of writing its output.

GFS has atomic rename that prevents output from being visible until complete.

Other failures/problems:

- For what applications **doesn't** MapReduce work well?

More flexible systems allow these, but more complex model.

Can use MapReduce for all these purposes!

```
reduce(cat id, list("dog!")) -> list(cat id)
```

Have fun with the lab!